

1.  $\lambda = 100 \text{ rps}$

$T_s = 0 - 19 \text{ ms}$  uniformly distributed

bound on F-Score unfairness

$$\frac{1}{2} \lambda T_q = \frac{1}{2} q \text{ by Little's Law}$$

To solve for  $q$  we must use M/G/1 formula

$$q = \frac{\rho^2 A}{1 - \rho} + \rho \quad A = \frac{1}{2} \left[ 1 + \left( \frac{\sigma T_s}{T_s} \right)^2 \right]$$

$$\rho = \frac{\lambda}{\mu} \quad \mu = \frac{1}{T_s}$$

$$T_s = \frac{0 + 19}{2} = 9.5 \text{ ms} = .0095 \text{ s}$$

$$\sigma T_s = \sqrt{\frac{1}{12} (b - a)^2} = \sqrt{\frac{1}{12} (19)^2} = 5.5 \text{ ms} = .0055 \text{ s}$$

$$A = \frac{1}{2} \left[ 1 + \left( \frac{.0055}{.0095} \right)^2 \right] = .67$$

$$\mu = \frac{1}{.0095} = 105$$

$$\rho = \frac{100}{105} = .95$$

$$q = \frac{.95^2 (.67)}{1 - .95} + .95 = 12.96$$

So the bound is  $(\frac{1}{2}) 12.96$  or  $6.5$



3.

a. The most obvious example would be if Q1 always has requests in it then requests in Q2 will never be serviced.

b. n - batch scan

c. A larger N makes the cache more efficient

d. A larger N is less fair though.

e.

□  $N=100$  The load on the system is low so we don't care about fairness. Setting N to 100 will also increase our systems performance.

□  $N=10$  The load is moderate so we should be more concerned about fairness.  $N=10$  provides a good balance on efficiency and fairness.

□  $N=1$  high load so fairness is a concern, caching provides little benefit here because there is little locality.

□  $N=10$  because the load is extremely high we should be concerned with fairness, however extremely high locality means caching will be effective, 10 provides a good balance between the two.



4.

a. yes starvation is possible. Class C jobs are the most susceptible because class B jobs are scheduled more frequently than class C jobs and have a higher priority.

b.

A event	$1/\text{minute}$	highest priority	5 seconds
B event	$1/\text{second}$	medium priority	.5 seconds
C event	$1/5 \text{ minutes}$	lowest priority	20 seconds

A worst case 5 seconds, it has the highest priority

B worst case 5.5 seconds. an A event arrives during the B events execution.

C worst case 50 seconds. I used geometric series to think about this

event C takes 20 seconds to do work which will generate 20 B events, Those 20 B events take 10s to service so they will generate an additional 10 B events those 10 B events take 5s to service so they will generate 5 more B events. This is a geometric series so the 20 seconds of work on A create 20 seconds of work on B events. If an A event arrives its 5 seconds of work will generate 5 seconds of work on B events. Total is  $20 + 20 + 5 + 5 = 50$ .



$$C \cdot 2X + \left( \left\lceil \frac{2X}{60} \right\rceil 5 \right) 2 = 300$$

$X = 125$  So 125 is the maximum amount of cpu time for C.

- d. if a Job C arrives just before a Job A the Job A will have to wait for Job C to finish using the resource R before it can begin to be serviced. Therefore the higher priority Job A must wait for the lower priority Job C.
- e. A Job C arrives before Job A. Then the C job is interrupted by Job B's causing Job C to take 40 seconds to complete Job A then runs for a total of 45 seconds.
- f. The above is an example of priority inversion even though Job A has higher priority than Job B. Job B's are serviced before Job A. A higher priority job is preempted by a lower priority one.
- g. Priority inheritance attempts to solve the problem of priority inversion. The priority of a low priority task is increased to the highest priority task waiting for the shared resource being consumed by the low priority task.
- h. Job B's can no longer interrupt Job C's using the resource so its worst case service time is 25 s